


QKSieve speedup decomposition across context length



Solid markers: same-protocol steady decode (2K-32K). Open markers: independent Llama measurements (64K-128K). Attention speedup and ceiling are latency-decomposition estimates.